

Lab7 - Introduction to linear regression

Michael Y.

April 14th, 2019

Batter up

The movie Moneyball focuses on the “quest for the secret of success in baseball”. It follows a low-budget team, the Oakland Athletics, who believed that underused statistics, such as a player’s ability to get on base, better predict the ability to score runs than typical statistics like home runs, RBIs (runs batted in), and batting average. Obtaining players who excelled in these underused statistics turned out to be much more affordable for the team.

In this lab we’ll be looking at data from all 30 Major League Baseball teams and examining the linear relationship between runs scored in a season and a number of other player statistics. Our aim will be to summarize these relationships both graphically and numerically in order to find which variable, if any, helps us best predict a team’s runs scored in a season.

The data

Let’s load up the data for the 2011 season.

```
load("more/mlb11.RData")
```

In addition to runs scored, there are seven traditionally used variables in the data set: at-bats, hits, home runs, batting average, strikeouts, stolen bases, and wins. There are also three newer variables: on-base percentage, slugging percentage, and on-base plus slugging. For the first portion of the analysis we’ll consider the seven traditional variables. At the end of the lab, you’ll work with the newer variables on your own.


1. What type of plot would you use to display the relationship between **runs** and one of the other numerical variables? Plot this relationship using the variable **at_bats** as the predictor. Does the relationship look linear? If you knew a team’s **at_bats**, would you be comfortable using a linear model to predict the number of runs?

Scatterplot.

Plot this relationship using the variable at_bats as the predictor.

```
plot(x = mlb11$at_bats, y=mlb11$runs)
title(main="Plot of at-bats vs. runs for 30 MLB teams in 2011")
```

Plot of at-bats vs. runs for 30 MLB teams in 2011



Does the relationship look linear?

It's difficult to discern visually whether or not there is a linear relationship because the fit is not very tight.

If you knew a team's at_bats, would you be comfortable using a linear model to predict the number of runs?

No, I don't believe that the relationship between these two variables provides a strong enough prediction to achieve the necessary level of comfort.

If the relationship looks linear, we can quantify the strength of the relationship with the correlation coefficient.

```
cor(mlb11$runs, mlb11$at_bats)
```

```
## [1] 0.61062705
```

Here the correlation is 0.61062705 which means that the goodness-of-fit, or R-Squared, is only 0.37286539 . This is not very strong.

Sum of squared residuals

Think back to the way that we described the distribution of a single variable. Recall that we discussed characteristics such as center, spread, and shape. It's also useful to be able to describe the relationship of


two numerical variables, such as `runs` and `at_bats` above.

2. Looking at your plot from the previous exercise, describe the relationship between these two variables. Make sure to discuss the form, direction, and strength of the relationship as well as any unusual observations.

The relationship appears to be somewhat linear, with an increasing relationship between at-bats and runs, but the relationship is not extremely strong.

Just as we used the mean and standard deviation to summarize a single variable, we can summarize the relationship between these two variables by finding the line that best follows their association. Use the following interactive function to select the line that you think does the best job of going through the cloud of points.

```
plot_ss(x = mlb11$at_bats, y = mlb11$runs)
```



```
## Click two points to make a line.
```

```
## Call:
```

```
## lm(formula = y ~ x, data = pts)
```

```
##
```

```
## Coefficients:
```


```
## (Intercept)          x
```

```
## -2789.24289      0.63055
```

```
##
```

```
## Sum of Squares: 123721.87
```

```
plot_ss(x = mlb11$at_bats, y = mlb11$runs
#       x1 = mean(mlb11$at_bats), y1=mean(mlb11$runs),
#       x2 = mean(mlb11$at_bats)+100, y2=mean(mlb11$runs)+100*.63055
#       )
```



```
## Click two points to make a line.
```

```
## Call:
```

```
## lm(formula = y ~ x, data = pts)
```

```
##
```

```
## Coefficients:
```

```
## (Intercept)          x
```

```
## -2789.24289      0.63055
```

```
##
```


```
## Sum of Squares: 123721.87
```

After running this command, you'll be prompted to click two points on the plot to define a line. Once you've done that, the line you specified will be shown in black and the residuals in blue. Note that there are 30 residuals, one for each of the 30 observations. Recall that the residuals are the difference between the observed values and the values predicted by the line:

$$e_i = y_i - \hat{y}_i$$

The most common way to do linear regression is to select the line that minimizes the sum of squared residuals. To visualize the squared residuals, you can rerun the plot command and add the argument `showSquares = TRUE`.

```
plot_ss(x = mlb11$at_bats, y = mlb11$runs, showSquares = TRUE)
```



```
## Click two points to make a line.
```

```
## Call:
```

```
## lm(formula = y ~ x, data = pts)
```

```
##
```

```
## Coefficients:
```

```
## (Intercept)          x
```

```
## -2789.24289      0.63055
```

```
##
```

```
## Sum of Squares: 123721.87
```

Note that the output from the `plot_ss` function provides you with the slope and intercept of your line as well as the sum of squares.


- Using `plot_ss`, choose a line that does a good job of minimizing the sum of squares. Run the function several times. What was the smallest sum of squares that you got? How does it compare to your neighbors?

The regression line must pass through the point $[\text{mean}(x), \text{mean}(y)]$.

Here, those figures are $[5523.5, 693.6]$

Knowing that the slope of the regression line is about $+0.63$, we can select another point by increasing x by 100 and y by 63. This gives a residual sum of squares equal to 123721.9. The best possible result is a tiny bit smaller:

```
plot_ss(x = mlb11$at_bats, y = mlb11$runs, showSquares = TRUE, leastSquares = T)
```



```
##
## Call:
## lm(formula = y ~ x)
##
## Coefficients:
## (Intercept)          x
## -2789.24289      0.63055
##
## Sum of Squares:  123721.87
```

The linear model

It is rather cumbersome to try to get the correct least squares line, i.e. the line that minimizes the sum of squared residuals, through trial and error. Instead we can use the `lm` function in R to fit the linear model (a.k.a. regression line).

```
m1 <- lm(runs ~ at_bats, data = mlb11)
```

The first argument in the function `lm` is a formula that takes the form $y \sim x$. Here it can be read that we want to make a linear model of `runs` as a function of `at_bats`. The second argument specifies that R should

look in the `mlb11` data frame to find the `runs` and `at_bats` variables.

The output of `lm` is an object that contains all of the information we need about the linear model that was just fit. We can access this information using the `summary` function.

```
summary(m1)

##
## Call:
## lm(formula = runs ~ at_bats, data = mlb11)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -125.576  -47.050  -16.588   54.399  176.868
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -2789.24289   853.69572  -3.2673 0.0028706 **
## at_bats      0.63055     0.15454   4.0801 0.0003388 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 66.473 on 28 degrees of freedom
## Multiple R-squared:  0.37287,    Adjusted R-squared:  0.35047
## F-statistic: 16.648 on 1 and 28 DF,  p-value: 0.00033884
```

Let's consider this output piece by piece. First, the formula used to describe the model is shown at the top. After the formula you find the five-number summary of the residuals. The "Coefficients" table shown next is key; its first column displays the linear model's y-intercept and the coefficient of `at_bats`. With this table, we can write down the least squares regression line for the linear model:

$$\hat{y} = -2789.2429 + 0.6305 * atbats$$

One last piece of information we will discuss from the summary output is the Multiple R-squared, or more simply, R^2 . The R^2 value represents the proportion of variability in the response variable that is explained by the explanatory variable. For this model, 37.3% of the variability in runs is explained by at-bats.

4. Fit a new model that uses `homeruns` to predict `runs`. Using the estimates from the R output, write the equation of the regression line. What does the slope tell us in the context of the relationship between success of a team and its home runs?

```
m2 <- lm(formula = runs ~ homeruns, data=mlb11)
summary(m2)

##
## Call:
## lm(formula = runs ~ homeruns, data = mlb11)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -91.6146 -33.4096   3.2308  24.2925 104.6306
##
## Coefficients:
##              Estimate Std. Error t value      Pr(>|t|)
## (Intercept)  415.23888   41.67789   9.9630 0.0000000001044 ***
## homeruns      1.83454    0.26765   6.8541 0.0000001900086 ***
```

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 51.295 on 28 degrees of freedom
## Multiple R-squared:  0.62656,    Adjusted R-squared:  0.61323
## F-statistic: 46.979 on 1 and 28 DF,  p-value: 0.00000019001
```

Using the estimates from the R output, write the equation of the regression line.


$$\hat{runs} = 415.23888 + 1.83454 * homeruns$$

What does the slope tell us in the context of the relationship between success of a team and its home runs?

The slope tells us that the number of runs scored by a team increases by 1.83454 with each homerun. Of course, this coefficient must be at least 1, because a solo HR (with no base runners) would score a single run (i.e., the batter), while homeruns hit with runners on base would result in 2, 3, or in the case of a “grand slam”, 4 runs.

```
plot(mlb11$runs ~ mlb11$homeruns)
abline(m2, col="blue", lwd=3)
title(main="Plot of homeruns vs. runs for 30 MLB teams in 2011 (slope = 1.83454)")
```

Plot of homeruns vs. runs for 30 MLB teams in 2011 (slope = 1.83454)




Prediction and prediction errors

Let's create a scatterplot with the least squares line laid on top.

```
plot(mlb11$runs ~ mlb11$at_bats)
abline(m1, col="red", lwd=3)
title(main="Plot of at-bats vs. runs for 30 MLB teams in 2011 (slope = 0.63)")
```

Plot of at-bats vs. runs for 30 MLB teams in 2011 (slope = 0.63)



The function `abline` plots a line based on its slope and intercept. Here, we used a shortcut by providing the model `m1`, which contains both parameter estimates. This line can be used to predict y at any value of x . When predictions are made for values of x that are beyond the range of the observed data, it is referred to as *extrapolation* and is not usually recommended. However, predictions made within the range of the data are more reliable. They're also used to compute the residuals.

5. If a team manager saw the least squares regression line and not the actual data, how many runs would he or she predict for a team with 5,578 at-bats? Is this an overestimate or an underestimate, and by how much? In other words, what is the residual for this prediction?

```
### what is the predicted number of runs for a team with 5578 at-bats?
pred_5578 <- m1$coefficients[1] + 5578 * m1$coefficients[2]
pred_5578
```

```
## (Intercept)
## 727.96497
```

```
### Is there any team with 5578 at-bats?
mlb11$at_bats==5578
```

```
## [1] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
```

```
## [23] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
sum(mlb11$at_bats==5578)

## [1] 0
mlb11[mlb11$at_bats==5578,]

## [1] team runs at_bats hits homeruns bat_avg strikeouts stol
## [11] new_slug new_obs
## <0 rows> (or 0-length row.names)
### No, there is no such team. Perhaps 5578 was a typo?

### What is the closest number of at-bats to 5578?
sort(mlb11$at_bats)

## [1] 5417 5421 5421 5421 5436 5436 5441 5447 5452 5486 5487 5502 5508 5509 5513 5518 5528 5532 5544
## [28] 5659 5672 5710
### OK, there is a team with 5579 at-bats. Perhaps this is the figure that was intended?
mlb11$at_bats==5579

## [1] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE TRUE
## [23] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
### Which team is this?
mlb11[mlb11$at_bats==5579,-12]

##
## team runs at_bats hits homeruns bat_avg strikeouts stolen_bases wins new_onbase
## 16 Philadelphia Phillies 713 5579 1409 153 0.253 1024 96 102 0.323
### OK, Philadelphia, which is tagged as the 16th row of the dataframe:
head(mlb11[order(mlb11$at_bats,decreasing = T),-12],10)

##
## team runs at_bats hits homeruns bat_avg strikeouts stolen_bases wins new_onbase
## 2 Boston Red Sox 875 5710 1600 203 0.280 1108 102 90 0.349
## 4 Kansas City Royals 730 5672 1560 129 0.275 1006 153 71 0.329
## 1 Texas Rangers 855 5659 1599 210 0.283 930 143 96 0.340
## 14 Cincinnati Reds 735 5612 1438 183 0.256 1250 97 79 0.326
## 6 New York Mets 718 5600 1477 108 0.264 1085 130 77 0.335
## 10 Houston Astros 615 5598 1442 95 0.258 1164 118 56 0.311
## 11 Baltimore Orioles 708 5585 1434 191 0.257 1120 81 69 0.316
## 16 Philadelphia Phillies 713 5579 1409 153 0.253 1024 96 102 0.323
## 3 Detroit Tigers 787 5563 1540 169 0.277 1143 49 95 0.340
## 20 Toronto Blue Jays 743 5559 1384 186 0.249 1184 131 81 0.317
### What is the predicted number of runs for a team with 5579 at-bats?
pred_5579 <- m1$coefficients[1] + 5579 * m1$coefficients[2]
pred_5579

## (Intercept)
## 728.59552
### How many runs did this team actually score?
actual_5579 <- mlb11[mlb11$at_bats==5579,'runs']
actual_5579

## [1] 713
```

```
### What is the residual?
residual_5579 <- actual_5579 - pred_5579
residual_5579
```

```
## (Intercept)
## -15.595525
```

The predicted number of runs for a team with 5,578 at-bats would be 727.96497461 .

However, there is no team in the data with exactly 5,578 at-bats. The closest is Philadelphia, which had 5,579.

(Perhaps 5,578 was a typo, and 5,579 was intended?)

If so, the predicted number of runs for a team with 5,579 at-bats would be 728.5955246 .


This is an *overestimate*, as the actual number of runs scored by this team was 713 .

Accordingly, the residual is negative: -15.5955246 .

The graph below illustrates that the team numbered 16 (Philadelphia) appears just below the regression line (near the center of the graph):

```
plot(y=mlb11$runs, x=mlb11$at_bats, col="lightblue", pch=19, cex=2)
abline(m1, col="red", lwd=3)
text(runs~at_bats, labels=rownames(mlb11), data=mlb11, cex=0.9, font=2)
title(main="Plot of at-bats vs. runs for MLB teams in 2011, with regression line")
```

Plot of at-bats vs. runs for MLB teams in 2011, with regression line



Interestingly, Philadelphia won the most games during the regular season:

```
head(mlb11[order(mlb11$wins,decreasing = T),-12],10)
```

| ## | team | runs | at_bats | hits | homeruns | bat_avg | strikeouts | stolen_bases | wins | new_onbase |
|-------|-----------------------|------|---------|------|----------|---------|------------|--------------|------|------------|
| ## 16 | Philadelphia Phillies | 713 | 5579 | 1409 | 153 | 0.253 | 1024 | 96 | 102 | 0.323 |
| ## 7 | New York Yankees | 867 | 5518 | 1452 | 222 | 0.263 | 1138 | 147 | 97 | 0.343 |
| ## 1 | Texas Rangers | 855 | 5659 | 1599 | 210 | 0.283 | 930 | 143 | 96 | 0.340 |
| ## 8 | Milwaukee Brewers | 721 | 5447 | 1422 | 185 | 0.261 | 1083 | 94 | 96 | 0.325 |
| ## 3 | Detroit Tigers | 787 | 5563 | 1540 | 169 | 0.277 | 1143 | 49 | 95 | 0.340 |
| ## 19 | Arizona Diamondbacks | 731 | 5421 | 1357 | 172 | 0.250 | 1249 | 133 | 94 | 0.322 |
| ## 25 | Tampa Bay Rays | 707 | 5436 | 1324 | 172 | 0.244 | 1193 | 155 | 91 | 0.322 |
| ## 2 | Boston Red Sox | 875 | 5710 | 1600 | 203 | 0.280 | 1108 | 102 | 90 | 0.349 |
| ## 5 | St. Louis Cardinals | 762 | 5532 | 1513 | 162 | 0.273 | 978 | 57 | 90 | 0.341 |
| ## 26 | Atlanta Braves | 641 | 5528 | 1345 | 173 | 0.243 | 1260 | 77 | 89 | 0.308 |

However, in postseason play, Philadelphia lost their NL Division Series to St. Louis, a Wild-Card team which would go on to win the 2011 World Series.

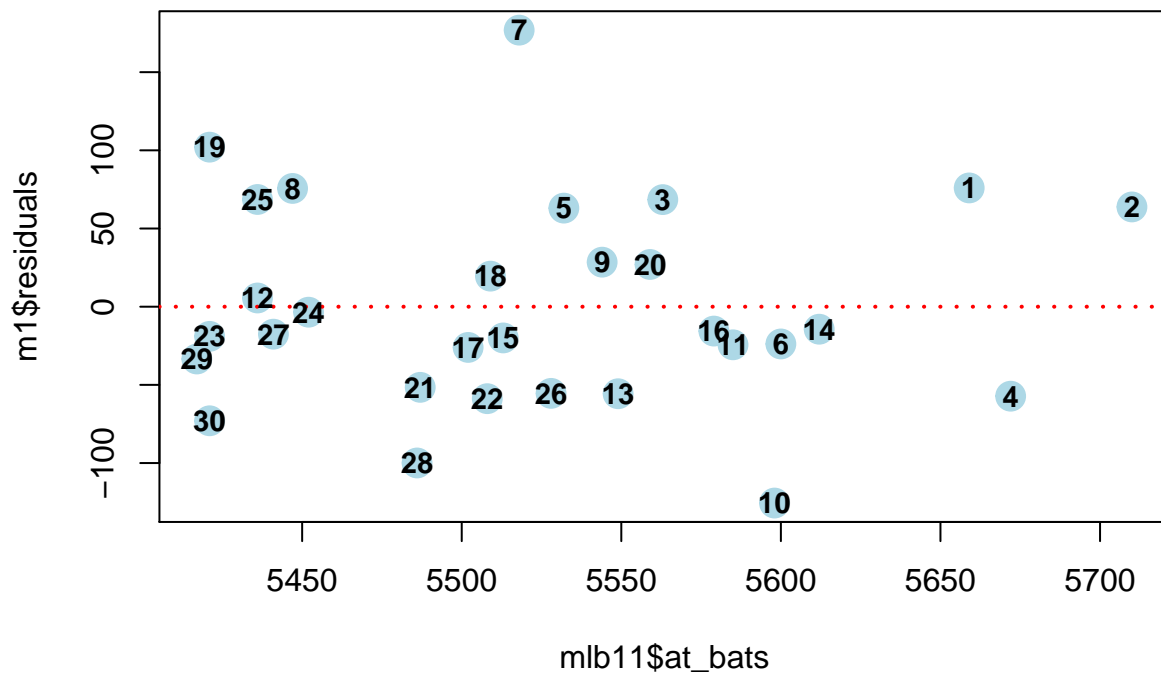
Model diagnostics

To assess whether the linear model is reliable, we need to check for (1) linearity, (2) nearly normal residuals, and (3) constant variability.

Linearity: You already checked if the relationship between runs and at-bats is linear using a scatterplot. We should also verify this condition with a plot of the residuals vs. at-bats. Recall that any code following a `#` is intended to be a comment that helps understand the code but is ignored by R.

```
plot(m1$residuals ~ mlb11$at_bats, col="lightblue", pch=19, cex=2)
abline(h = 0, lty = 3, col="red", lwd=2) # adds a horizontal dashed line at y = 0
text(m1$residuals~mlb11$at_bats, labels=rownames(mlb11),data=mlb11, cex=0.9, font=2)
title(main="Plot of residual of predicted runs vs. at-bats for MLB teams in 2011")
```

Plot of residual of predicted runs vs. at-bats for MLB teams in 2011



6. Is there any apparent pattern in the residuals plot? What does this indicate about the linearity of the relationship between runs and at-bats?

No, there does not appear to be any pattern in the residuals plot.

However, I note that there are more teams “below the line” (i.e., where runs have been over-predicted) vs. teams “above the line” (where teams have “exceeded expectations”).

Notably, success of team #7 (New York Yankees), for which their residual is so large (176.868, which is 2.66 standard deviations from the expectation) is an outlier which may have imposed unusual influence on the regression. Removal of this data point may provide more reasonable expectations for the other teams.

Using a linear model for the relationship between runs and at-bats may be appropriate.

To check for linearity, we can use the “modelAssumptions” test from the lmSupport package:

```
require(lmSupport)    ## Note: the "S" is capitalized in the package name
```

```
## Loading required package: lmSupport
```

```
## Warning: package 'lmSupport' was built under R version 3.5.3
```

```
modelAssumptions(m1,"LINEAR")
```

```
##
```

```
## Call:
```

```
## lm(formula = runs ~ at_bats, data = mlb11)
```

```
##
```

```
## Coefficients:
```

```
## (Intercept)      at_bats
```

```
## -2789.24289      0.63055
```

```
##
```

```
##
```

```
## ASSESSMENT OF THE LINEAR MODEL ASSUMPTIONS
```

```
## USING THE GLOBAL TEST ON 4 DEGREES-OF-FREEDOM:
```

```
## Level of Significance = 0.05
```

```
##
```

```
## Call:
```

```
## gvlma(x = Model)
```

```
##
```

```
##
```


| | Value | p-value | Decision |
|-----------------------|----------|---------|-------------------------|
| ## Global Stat | 3.361222 | 0.49929 | Assumptions acceptable. |
| ## Skewness | 1.479243 | 0.22389 | Assumptions acceptable. |
| ## Kurtosis | 0.079467 | 0.77802 | Assumptions acceptable. |
| ## Link Function | 0.508898 | 0.47562 | Assumptions acceptable. |
| ## Heteroscedasticity | 1.293614 | 0.25538 | Assumptions acceptable. |

The test results are consistent with Linearity.

Nearly normal residuals: To check this condition, we can look at a histogram:


```
hist(m1$residuals)
```

Histogram of m1\$residuals



or a normal probability plot of the residuals:

```
qqnorm(m1$residuals)
qqline(m1$residuals) # adds diagonal line to the normal prob plot
```



7. Based on the histogram and the normal probability plot, does the nearly normal residuals condition appear to be met?

The histogram does not appear to be “nearly normal” as the number of observations below zero (18) is greater than the number above zero (12). Likewise, the results on Q-Q-plot indicate that the empirical quantiles are not as close to the normal quantiles as would be expected.

These results would suggest that the “Nearly-normal residuals” condition may not appear to be met. However, it would be better to perform an actual test for normality, such as Shapiro-Wilks:

```
shapiro.test(m1$residuals)

##
##  Shapiro-Wilk normality test
##
## data:  m1$residuals
## W = 0.961442, p-value = 0.33703
```

Because the p-value is large, we *fail to reject* the Null Hypothesis, which is that the residuals *ARE* normal.

Another useful test for normality is Kolmogorov-Smirnov. Here we test whether the residuals are consistent with a Normal distribution with mean 0 and with standard deviation matching

that of the residuals:

```
ks.test(m1$residuals,"pnorm",0,sd(m1$residuals))
```

```
##  
## One-sample Kolmogorov-Smirnov test  
##  
## data: m1$residuals  
## D = 0.153934, p-value = 0.43254  
## alternative hypothesis: two-sided
```

Here as well, the large p-value indicates that we *fail to reject* the null hypothesis, which is that the residuals are normal.

Another useful test of normality is the Anderson Darling test:

```
require(nortest)
```

```
## Loading required package: nortest
```

```
ad.test(m1$residuals)
```

```
##  
## Anderson-Darling normality test  
##  
## data: m1$residuals  
## A = 0.515229, p-value = 0.17651
```

Here again, the high p-value indicates that we *fail to reject* the null hypothesis, which is that the residuals are normal.

Yet another useful test for normality is the Jarque-Bera test:

```
require(tseries)
```

```
## Loading required package: tseries
```

```
## Warning: package 'tseries' was built under R version 3.5.3
```

```
jarque.bera.test(m1$residuals)
```

```
##  
## Jarque Bera Test  
##  
## data: m1$residuals  
## X-squared = 1.55871, df = 2, p-value = 0.4587
```

Here again, the high p-value indicates that we *fail to reject* the null hypothesis, which is that the residuals are normal.

Despite the questionable nature of the histogram and the QQ-plot, these numerical tests of the residuals fail to reject normality. This provides evidence that the residuals *are* normal.

Constant variability:

8. Based on the plot in (1), does the constant variability condition appear to be met?

Yes, the “constant variability” (i.e., homoscedasticity) condition does appear to be met.

A useful numeric test for constant variance is Breusch-Pagan. As it assumes that the data are normally distributed, the above tests need to have passed before we can use it.

```
require(olsrr)

## Loading required package: olsrr
## Warning: package 'olsrr' was built under R version 3.5.3
##
## Attaching package: 'olsrr'
## The following object is masked from 'package:datasets':
##
##     rivers

ols_test_breusch_pagan(m1)

##
## Breusch Pagan Test for Heteroskedasticity
## -----
## Ho: the variance is constant
## Ha: the variance is not constant
##
##           Data
## -----
## Response : runs
## Variables: fitted values of runs
##
##           Test Summary
## -----
## DF          =    1
## Chi2         =   0.014290331
## Prob > Chi2  =   0.90484583
```

The high p-value indicates that we fail to reject H_0 , which is that the variance is constant.

```
allnames = names(mlb11)

tradvars = allnames[3:9]
numtradvars = length(tradvars) # 7
```

On Your Own

9. Choose any traditional variable from `mlb11` that you think might be a good predictor of `runs`. Produce a scatterplot of the two variables and fit a linear model. At a glance, does there seem to be a linear relationship?

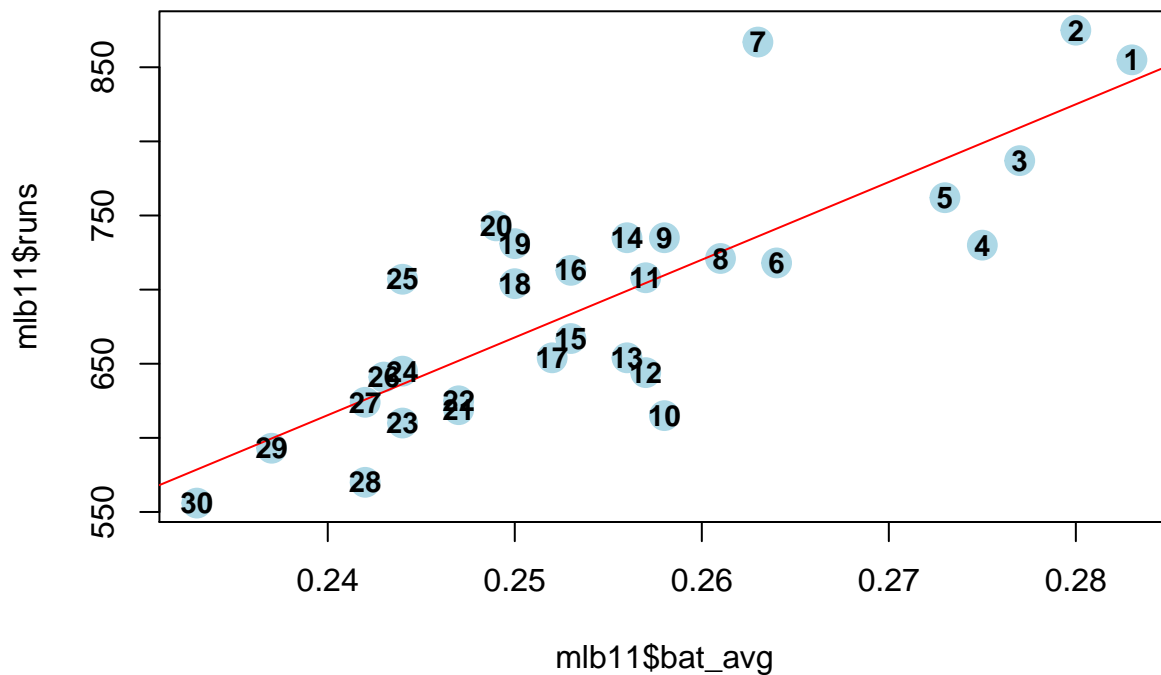
The “traditional” variables include the following:

```
[1] "at_bats"
[2] "hits"
[3] "homeruns"
[4] "bat_avg"
[5] "strikeouts"
[6] "stolen_bases"
[7] "wins"
```

I choose “bat_avg” – the team batting average.

```
plot(x = mlb11$bat_avg , y = mlb11$runs, col="lightblue", pch=19, cex=2)
batavgmod <- lm(runs ~ bat_avg, data=mlb11)
abline(batavgmod, col="red")
text(runs~bat_avg, labels=rownames(mlb11),data=mlb11, cex=0.9, font=2)
title(main = "Plot of batting average vs. runs scored in 2011 by 30 MLB teams")
```

Plot of batting average vs. runs scored in 2011 by 30 MLB teams



10. How does this relationship compare to the relationship between `runs` and `at_bats`? Use the R^2 values from the two model summaries to compare. Does your variable seem to predict `runs` better than `at_bats`? How can you tell?

```
print("\n\n*****BATTING AVERAGE MODEL: ")
```

```
## [1] "\n\n*****BATTING AVERAGE MODEL: "
```

```
batavgmodsumm = summary(batavgmod)
print(batavgmodsumm)
```

```
##
## Call:
## lm(formula = runs ~ bat_avg, data = mlb11)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -94.6762 -26.3034  -5.4961  28.4822 131.1127
##
## Coefficients:
##              Estimate Std. Error t value    Pr(>|t|)
## (Intercept)  -642.82     183.08  -3.5111    0.001531 **
## bat_avg       5242.23     717.28   7.3085  0.00000005877 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 49.226 on 28 degrees of freedom
## Multiple R-squared:  0.65608,    Adjusted R-squared:  0.64379
## F-statistic: 53.414 on 1 and 28 DF,  p-value: 0.00000005877
```

```
batavg_R2 <- batavgmodsumm$r.squared
batavg_AdjR2 <- batavgmodsumm$adj.r.squared
```

```
print("\n\n*****AT-BATS MODEL: ")
```

```
## [1] "\n\n*****AT-BATS MODEL: "
```

```
atbatsmodsumm = summary(m1)
print(atbatsmodsumm)
```

```
##
## Call:
## lm(formula = runs ~ at_bats, data = mlb11)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -125.576  -47.050  -16.588   54.399  176.868
##
## Coefficients:
##              Estimate Std. Error t value    Pr(>|t|)
## (Intercept) -2789.24289    853.69572  -3.2673  0.0028706 **
## at_bats       0.63055      0.15454   4.0801  0.0003388 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 66.473 on 28 degrees of freedom
## Multiple R-squared:  0.37287,    Adjusted R-squared:  0.35047
## F-statistic: 16.648 on 1 and 28 DF,  p-value: 0.00033884
```

```
atbats_R2 <- atbatsmodsumm$r.squared
atbats_AdjR2 <- atbatsmodsumm$adj.r.squared
```

```
print(paste("Batting Average Model as predictor of runs: ", "R2: ", batavg_R2, "AdjR2: ", batavg_AdjR2
```

```
## [1] "Batting Average Model as predictor of runs: R2: 0.656077134646863 AdjR2: 0.643794175169965"
```

```
print(paste("At-Bats Model          as predictor of runs: ", "R2: ", atbats_R2, "AdjR2: ", atbats_AdjR2))

## [1] "At-Bats Model          as predictor of runs: R2:  0.372865390186805 AdjR2:  0.350467725550619"
```

How does this relationship compare to the relationship between runs and at_bats?

Because the R^2 and the $AdjustedR^2$ values for the Batting Average model are much greater than those of the At-Bats model, the Batting Average Model provides a much stronger fit and is preferred as a predictor.

***Does your variable seem to predict runs better than at_bats? How can you tell?

```
batavganova <- anova(batavgmod)
print(batavganova)

## Analysis of Variance Table
##
## Response: runs
##          Df    Sum Sq Mean Sq F value    Pr(>F)
## bat_avg    1 129431.7 129431.7  53.4136 0.00000005877 ***
## Residuals 28  67849.5   2423.2
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

batavg_ss <- batavganova$`Sum Sq`
print(batavg_ss)

## [1] 129431.684  67849.516

batvag_Explained_Sum_of_Squares <- batavg_ss[1]
batvag_Residual_Sum_of_Squares <- batavg_ss[2]

print(paste("batvag_Explained_Sum_of_Squares: ", batvag_Explained_Sum_of_Squares ))

## [1] "batvag_Explained_Sum_of_Squares:  129431.684415694"
print(paste("batvag_Residual_Sum_of_Squares: ", batvag_Residual_Sum_of_Squares ))

## [1] "batvag_Residual_Sum_of_Squares:  67849.5155843053"

atbatsanova <- anova(m1)
print(atbatsanova)

## Analysis of Variance Table
##
## Response: runs
##          Df    Sum Sq Mean Sq F value    Pr(>F)
## at_bats    1  73559.3  73559.3  16.6475 0.00033884 ***
## Residuals 28 123721.9   4418.6
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

atbats_ss <- atbatsanova$`Sum Sq`
print(atbats_ss)

## [1]  73559.332 123721.868
```

```

atbats_Explained_Sum_of_Squares <- atbats_ss[1]
atbats_Residual_Sum_of_Squares <- atbats_ss[2]

print(paste("atbats_Explained_Sum_of_Squares: ", atbats_Explained_Sum_of_Squares ))

## [1] "atbats_Explained_Sum_of_Squares: 73559.3316145209"
print(paste("atbats_Residual_Sum_of_Squares: ", atbats_Residual_Sum_of_Squares ))

## [1] "atbats_Residual_Sum_of_Squares: 123721.868385479"

```

The atbats Explained Sum of Squares is quite low, while the Residual Sum of Squares is high. This indicates that the model explains less than half of the variance.

In contrast, the batting average Explained Sum of Squared is quite high, while the residual Sum of Squares is low. This indicates that this model explains much more of the variance, and thus provides predictions which have smaller errors than the atbats model.

11. Now that you can summarize the linear relationship between two variables, investigate the relationships between runs and each of the other five traditional variables. Which variable best predicts runs? Support your conclusion using the graphical and numerical methods we've discussed (for the sake of conciseness, only include output for the best variable, not all five).

Evaluating the single-variable model for each of the seven “Traditional variables:”

```

allnames = names(mlb11)

tradvars = allnames[3:9]
numtradvars = length(tradvars) # 7

#newvars = allnames[10:12]
#numnewvars = length(newvars) # 3

#allvars = allnames[3:12]
#numallvars = length(allvars) # 10

trad_r2 = array(data = 0, dim=numtradvars, dimnames=list(tradvars))
trad_adj_r2 = array(data = 0, dim=numtradvars, dimnames=list(tradvars))

for (i in 1:numtradvars) {
  name = tradvars[i]
  print(paste ( "\n\n*****Doing ", name, "..."))
  fmula = paste("runs ~", name)
  print(paste ( "Formula: ", fmula))
  mod = lm(fmula, data=mlb11)
  print(      "Model: ")
  print(mod)
  summod = summary(mod)
  print(      "Summod: ")
  print(summod)
  trad_r2[i]=summod$r.squared
  print(paste ( "R2: ", trad_r2[i]))
  trad_adj_r2[i] = summod$adj.r.squared
}

```

```

print(paste ( "AdjR2: ", trad_adjr2[i]))

print(paste ("***** DONE WITH ", name, "*****"))
}

## [1] "\n\n*****Doing at_bats ..."
## [1] "Formula: runs ~ at_bats"
## [1] "Model: "
##
## Call:
## lm(formula = fmula, data = mlb11)
##
## Coefficients:
## (Intercept) at_bats
## -2789.24289 0.63055
##
## [1] "Summod: "
##
## Call:
## lm(formula = fmula, data = mlb11)
##
## Residuals:
## Min 1Q Median 3Q Max
## -125.576 -47.050 -16.588 54.399 176.868
##
## Coefficients:
## Estimate Std. Error t value Pr(>|t|)
## (Intercept) -2789.24289 853.69572 -3.2673 0.0028706 **
## at_bats 0.63055 0.15454 4.0801 0.0003388 ***
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 66.473 on 28 degrees of freedom
## Multiple R-squared: 0.37287, Adjusted R-squared: 0.35047
## F-statistic: 16.648 on 1 and 28 DF, p-value: 0.00033884
##
## [1] "R2: 0.372865390186805"
## [1] "AdjR2: 0.350467725550619"
## [1] "***** DONE WITH at_bats *****"
## [1] "\n\n*****Doing hits ..."
## [1] "Formula: runs ~ hits"
## [1] "Model: "
##
## Call:
## lm(formula = fmula, data = mlb11)
##
## Coefficients:
## (Intercept) hits
## -375.55997 0.75886
##
## [1] "Summod: "
##
## Call:
## lm(formula = fmula, data = mlb11)

```

```

##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -103.7183  -27.1794   -5.2328   19.3224  140.6931
##
## Coefficients:
##              Estimate Std. Error t value    Pr(>|t|)
## (Intercept) -375.55997   151.18056  -2.4842    0.01924 *
## hits         0.75886     0.10711   7.0851 0.0000001043 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 50.228 on 28 degrees of freedom
## Multiple R-squared:  0.64194,    Adjusted R-squared:  0.62915
## F-statistic: 50.199 on 1 and 28 DF,  p-value: 0.00000010432
##
## [1] "R2: 0.641938767239419"
## [1] "AdjR2: 0.629150866069399"
## [1] "***** DONE WITH hits *****"
## [1] "\n\n*****Doing homeruns ..."
## [1] "Formula: runs ~ homeruns"
## [1] "Model: "
##
## Call:
## lm(formula = fmula, data = mlb11)
##
## Coefficients:
## (Intercept)      homeruns
##    415.2389         1.8345
##
## [1] "Summod: "
##
## Call:
## lm(formula = fmula, data = mlb11)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -91.6146 -33.4096   3.2308  24.2925 104.6306
##
## Coefficients:
##              Estimate Std. Error t value    Pr(>|t|)
## (Intercept) 415.23888   41.67789   9.9630 0.0000000001044 ***
## homeruns     1.83454     0.26765   6.8541 0.0000001900086 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 51.295 on 28 degrees of freedom
## Multiple R-squared:  0.62656,    Adjusted R-squared:  0.61323
## F-statistic: 46.979 on 1 and 28 DF,  p-value: 0.00000019001
##
## [1] "R2: 0.626563569566283"
## [1] "AdjR2: 0.61322655419365"
## [1] "***** DONE WITH homeruns *****"
## [1] "\n\n*****Doing bat_avg ..."

```



```

## [1] "Formula: runs ~ bat_avg"
## [1] "Model: "
##
## Call:
## lm(formula = fmula, data = mlb11)
##
## Coefficients:
## (Intercept)      bat_avg
##      -642.82      5242.23
##
## [1] "Summod: "
##
## Call:
## lm(formula = fmula, data = mlb11)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -94.6762 -26.3034  -5.4961  28.4822 131.1127
##
## Coefficients:
##              Estimate Std. Error t value    Pr(>|t|)
## (Intercept)  -642.82     183.08  -3.5111    0.001531 **
## bat_avg      5242.23     717.28   7.3085  0.00000005877 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 49.226 on 28 degrees of freedom
## Multiple R-squared:  0.65608,    Adjusted R-squared:  0.64379
## F-statistic: 53.414 on 1 and 28 DF,  p-value: 0.00000005877
##
## [1] "R2: 0.656077134646863"
## [1] "AdjR2: 0.643794175169965"
## [1] "***** DONE WITH bat_avg *****"
## [1] "\n\n*****Doing strikeouts ..."
## [1] "Formula: runs ~ strikeouts"
## [1] "Model: "
##
## Call:
## lm(formula = fmula, data = mlb11)
##
## Coefficients:
## (Intercept) strikeouts
##  1054.73423    -0.31414
##
## [1] "Summod: "
##
## Call:
## lm(formula = fmula, data = mlb11)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -132.270  -46.948  -11.919   55.137  169.756
##
## Coefficients:

```

```

##               Estimate Std. Error t value      Pr(>|t|)
## (Intercept) 1054.73423  151.78899  6.9487 0.0000001486 ***
## strikeouts   -0.31414    0.13148 -2.3893   0.02386 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 76.502 on 28 degrees of freedom
## Multiple R-squared:  0.16936,    Adjusted R-squared:  0.13969
## F-statistic: 5.7089 on 1 and 28 DF,  p-value: 0.023856
##
## [1] "R2: 0.169357932236313"
## [1] "AdjR2: 0.139692144101895"
## [1] "***** DONE WITH strikeouts *****"
## [1] "\n\n*****Doing stolen_bases ..."
## [1] "Formula: runs ~ stolen_bases"
## [1] "Model: "
##
## Call:
## lm(formula = fmula, data = mlb11)
##
## Coefficients:
## (Intercept) stolen_bases
##      677.30743      0.14906
##
## [1] "Summod: "
##
## Call:
## lm(formula = fmula, data = mlb11)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -139.940  -62.874   10.007   38.544  182.488
##
## Coefficients:
##               Estimate Std. Error t value      Pr(>|t|)
## (Intercept)  677.30743   58.97508 11.4846 0.000000000004166 ***
## stolen_bases   0.14906    0.52109  0.2861    0.7769
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 83.817 on 28 degrees of freedom
## Multiple R-squared:  0.002914,    Adjusted R-squared: -0.032696
## F-statistic: 0.08183 on 1 and 28 DF,  p-value: 0.77694
##
## [1] "R2: 0.00291399266657394"
## [1] "AdjR2: -0.0326962218810485"
## [1] "***** DONE WITH stolen_bases *****"
## [1] "\n\n*****Doing wins ..."
## [1] "Formula: runs ~ wins"
## [1] "Model: "
##
## Call:
## lm(formula = fmula, data = mlb11)
##

```

```
## Coefficients:
## (Intercept)      wins
##      342.121      4.341
##
## [1] "Summod: "
##
## Call:
## lm(formula = fmula, data = mlb11)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -145.4498  -47.5062   -7.4819   47.3463  142.1860
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  342.1214     89.2230   3.8345 0.0006538 ***
## wins          4.3410       1.0915   3.9770 0.0004469 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 67.1 on 28 degrees of freedom
## Multiple R-squared:  0.36097,    Adjusted R-squared:  0.33815
## F-statistic: 15.816 on 1 and 28 DF,  p-value: 0.00044693
##
## [1] "R2: 0.360971179446681"
## [1] "AdjR2: 0.338148721569776"
## [1] "***** DONE WITH wins *****"
```

R^2 and *Adjusted* – R^2 for single-predictor models containing each of the “traditional” variables:

```
print("R-SQUARED: ")

## [1] "R-SQUARED: "
print(t(t(sort(trad_r2,decreasing = T))))

##              [,1]
## bat_avg      0.6560771346
## hits         0.6419387672
## homeruns     0.6265635696
## at_bats      0.3728653902
## wins         0.3609711794
## strikeouts   0.1693579322
## stolen_bases 0.0029139927
print("Adjusted R-SQUARED: ")

## [1] "Adjusted R-SQUARED: "
print(t(t(sort(trad_adjr2,decreasing = T))))

##              [,1]
## bat_avg      0.643794175
## hits         0.629150866
## homeruns     0.613226554
```

```
## at_bats      0.350467726
## wins        0.338148722
## strikeouts   0.139692144
## stolen_bases -0.032696222
```

Which variable best predicts runs?

Among the “Traditional” variables, Batting Average, Hits, and HomeRuns each provide an R^2 and an *Adjusted – R^2* in the .60 range, which is rather strong, while At Bats and Wins provide results in the .30 range, which is only moderately strong. Strikeouts and Stolen Bases provide extremely weak predictions of runs.

The very best predictor is Batting Average . The graphical results for the model containing this variable are shown above in response to the question numbered (9) .

12. Now examine the three newer variables. These are the statistics used by the author of *Moneyball* to predict a teams success. In general, are they more or less effective at predicting runs than the old variables? Explain using appropriate graphical and numerical evidence. Of all ten variables we’ve analyzed, which seems to be the best predictor of runs? Using the limited (or not so limited) information you know about these baseball statistics, does your result make sense?

Evaluating the single-variable model for each of the three “New” variables:

```
allnames = names(mlb11)

#tradvars = allnames[3:9]
#numtradvars = length(tradvars) # 7

newvars = allnames[10:12]
numnewvars = length(newvars) # 3

#allvars = allnames[3:12]
#numallvars = length(allvars) # 10

new_r2 = array(data = 0, dim=numnewvars, dimnames=list(newvars))
new_adjR2 = array(data = 0, dim=numnewvars, dimnames=list(newvars))

for (i in 1:numnewvars) {
  name = newvars[i]
  print(paste ( "\n\n*****Doing ", name, "..."))
  fmula = paste("runs ~", name)
  print(paste ( "Formula: ", fmula))
  mod = lm(fmula, data=mlb11)
  print(      "Model: ")
  print(mod)
  summod = summary(mod)
  print(      "Summod: ")
  print(summod)
  new_r2[i]=summod$r.squared
  print(paste ( "R2: ", new_r2[i]))
  new_adjR2[i] = summod$adj.r.squared
  print(paste ( "AdjR2: ", new_adjR2[i]))
}
```

```

print(paste ("***** DONE WITH ", name, "*****"))
}

## [1] "\n\n*****Doing new_onbase ..."
## [1] "Formula: runs ~ new_onbase"
## [1] "Model: "
##
## Call:
## lm(formula = fmula, data = mlb11)
##
## Coefficients:
## (Intercept) new_onbase
## -1118.4 5654.3
##
## [1] "Summod: "
##
## Call:
## lm(formula = fmula, data = mlb11)
##
## Residuals:
## Min 1Q Median 3Q Max
## -58.2700 -18.3348 3.2486 19.5203 69.0016
##
## Coefficients:
## Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1118.42 144.48 -7.741 0.0000000196789595 ***
## new_onbase 5654.32 450.46 12.552 0.0000000000005116 ***
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 32.606 on 28 degrees of freedom
## Multiple R-squared: 0.84911, Adjusted R-squared: 0.84372
## F-statistic: 157.56 on 1 and 28 DF, p-value: 0.00000000000051157
##
## [1] "R2: 0.849105251446139"
## [1] "AdjR2: 0.843716153283501"
## [1] "***** DONE WITH new_onbase *****"
## [1] "\n\n*****Doing new_slug ..."
## [1] "Formula: runs ~ new_slug"
## [1] "Model: "
##
## Call:
## lm(formula = fmula, data = mlb11)
##
## Coefficients:
## (Intercept) new_slug
## -375.8 2681.3
##
## [1] "Summod: "
##
## Call:
## lm(formula = fmula, data = mlb11)
##
## Residuals:

```

```

##      Min      1Q   Median      3Q      Max
## -45.4096 -18.6566 -0.9096  16.2935  52.2932
##
## Coefficients:
##           Estimate Std. Error t value      Pr(>|t|)
## (Intercept) -375.804      68.708 -5.4696 0.00000769547969384 ***
## new_slug    2681.331     171.830 15.6046 0.000000000000000242 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 26.956 on 28 degrees of freedom
## Multiple R-squared:  0.89687,    Adjusted R-squared:  0.89319
## F-statistic: 243.5 on 1 and 28 DF,  p-value: 0.00000000000000024201
##
## [1] "R2: 0.896870368409638"
## [1] "AdjR2: 0.893187167281411"
## [1] "***** DONE WITH new_slug *****"
## [1] "\n\n*****Doing new_obs ..."
## [1] "Formula: runs ~ new_obs"
## [1] "Model: "
##
## Call:
## lm(formula = fmula, data = mlb11)
##
## Coefficients:
## (Intercept)      new_obs
##      -686.61      1919.36
##
## [1] "Summod: "
##
## Call:
## lm(formula = fmula, data = mlb11)
##
## Residuals:
##      Min      1Q   Median      3Q      Max
## -43.4563 -13.6903   1.1646  13.9352  41.1559
##
## Coefficients:
##           Estimate Std. Error t value      Pr(>|t|)
## (Intercept) -686.614      68.925 -9.9617 0.0000000001047 ***
## new_obs     1919.364      95.695 20.0571 < 0.00000000000000022 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 21.412 on 28 degrees of freedom
## Multiple R-squared:  0.93493,    Adjusted R-squared:  0.9326
## F-statistic: 402.29 on 1 and 28 DF,  p-value: < 0.000000000000000222
##
## [1] "R2: 0.934927126351814"
## [1] "AdjR2: 0.932603095150093"
## [1] "***** DONE WITH new_obs *****"

```

R^2 and *Adjusted* $-R^2$ for single-predictor models containing each of the three “new” variables:

```
print("R-SQUARED: ")

## [1] "R-SQUARED: "
print(t(t(sort(new_r2,decreasing = T))))

##           [,1]
## new_obs    0.93492713
## new_slug   0.89687037
## new_onbase 0.84910525
print("Adjusted R-SQUARED: ")

## [1] "Adjusted R-SQUARED: "
print(t(t(sort(new_adjr2,decreasing = T))))

##           [,1]
## new_obs    0.93260310
## new_slug   0.89318717
## new_onbase 0.84371615
```

In general, are they more or less effective at predicting runs than the old variables? Explain using appropriate graphical and numerical evidence.

Because the “New” variables have much higher R^2 than the “Traditional” variables, they would be more effective at predicting runs. The R^2 and *Adjusted* R^2 values for each variable are given above.

Of all ten variables we’ve analyzed, which seems to be the best predictor of runs?

The variable new_obs (“on-base plus slugging”) seems to be the best predictor, as it has the highest R^2 .

```
new_obs_model <- lm(formula = runs ~ new_obs, data = mlb11)
summary(new_obs_model)

##
## Call:
## lm(formula = runs ~ new_obs, data = mlb11)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -43.4563 -13.6903  1.1646  13.9352  41.1559
##
## Coefficients:
##              Estimate Std. Error t value      Pr(>|t|)
## (Intercept)  -686.614     68.925  -9.9617  0.0000000001047 ***
## new_obs       1919.364     95.695  20.0571 < 0.0000000000000022 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 21.412 on 28 degrees of freedom
```

```
## Multiple R-squared:  0.93493,    Adjusted R-squared:  0.9326
## F-statistic: 402.29 on 1 and 28 DF,  p-value: < 0.000000000000000222
```

```
anova(new_obs_model)
```

```
## Analysis of Variance Table
```

```
##
```

```
## Response: runs
```

```
##           Df    Sum Sq Mean Sq F value           Pr(>F)
```

```
## new_obs     1 184443.5 184443.5 402.287 < 0.000000000000000222 ***
```

```
## Residuals  28  12837.7    458.5
```

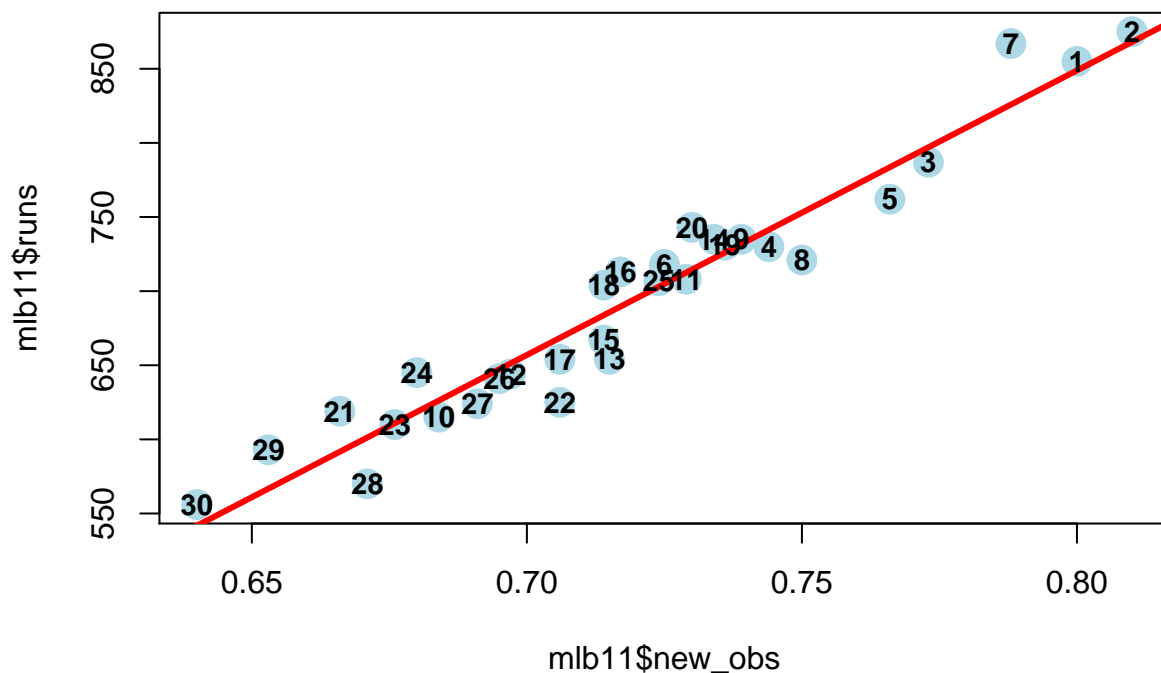
```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

As can be seen above, the R^2 results are above 90 percent, and the explained sum of squares (184443.5) far exceeds the Residual Sum of Squares (12837.7). This indicates an extremely tight fit.

```
plot(y=mlb11$runs, x=mlb11$new_obs, col="lightblue", pch=19, cex=2)
abline(new_obs_model, col="red", lwd=3)
text(runs~new_obs, labels=rownames(mlb11), data=mlb11, cex=0.9, font=2)
title(main="Plot of new_obs (On Base plus Slugging) vs. runs for MLB teams in 2011,\nwith regression line")
```

Plot of new_obs (On Base plus Slugging) vs. runs for MLB teams in 2011 with regression line



The above plot confirms an extremely tight fit, which makes for accurate predictions.

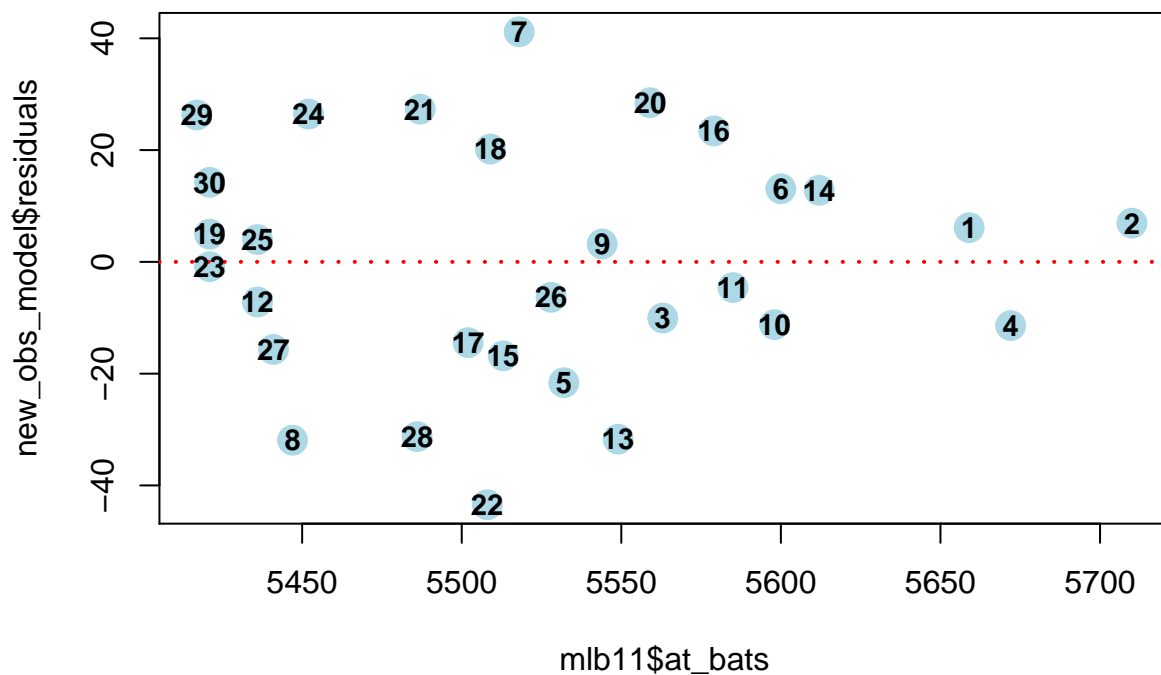
13. Check the model diagnostics for the regression model with the variable you decided was the best predictor for runs.

To assess whether the linear model is reliable, we need to check for

- (1) linearity,
- (2) nearly normal residuals, and
- (3) constant variability.

```
plot(new_obs_model$residuals ~ mlb11$at_bats, col="lightblue", pch=19, cex=2)
abline(h = 0, lty = 3, col="red", lwd=2) # adds a horizontal dashed line at y = 0
text(new_obs_model$residuals~mlb11$at_bats, labels=rownames(mlb11),data=mlb11, cex=0.9, font=2)
title(main="Plot of residual of predicted runs vs. OBS (On-Base plus Slugging) \nfor MLB teams in 2011")
```

**Plot of residual of predicted runs vs. OBS (On-Base plus Slugging)
for MLB teams in 2011**



To check for linearity, we can use the “modelAssumptions” test from the lmSupport package:

```
require(lmSupport) ## Note: the "S" is capitalized in the package name
modelAssumptions(new_obs_model,"LINEAR")
```


```
##
## Call:
## lm(formula = runs ~ new_obs, data = mlb11)
##
## Coefficients:
## (Intercept)      new_obs
##      -686.61      1919.36
##
##
```

```
## ASSESSMENT OF THE LINEAR MODEL ASSUMPTIONS
## USING THE GLOBAL TEST ON 4 DEGREES-OF-FREEDOM:
## Level of Significance = 0.05
##
## Call:
## gvlma(x = Model)
##
##               Value p-value              Decision
## Global Stat    3.711018 0.44653 Assumptions acceptable.
## Skewness       0.042733 0.83623 Assumptions acceptable.
## Kurtosis       0.610288 0.43468 Assumptions acceptable.
## Link Function   2.615006 0.10586 Assumptions acceptable.
## Heteroscedasticity 0.442990 0.50568 Assumptions acceptable.
```

The test results are consistent with Linearity.


Nearly normal residuals: To check this condition, we can look at a histogram:

```
hist(new_obs_model$residuals)
```



or a normal probability plot of the residuals:

```
qqnorm(new_obs_model$residuals)
qqline(new_obs_model$residuals) # adds diagonal line to the normal prob plot
```



The above plots confirm that the residuals from the OBS model are “nearly normal.”

```
### Shapiro-Wilks test of Normality:
shapiro.test(new_obs_model$residuals)
```

```
##
##  Shapiro-Wilk normality test
##
## data:  new_obs_model$residuals
## W = 0.980683, p-value = 0.84342
```

```
### Kolmogorov-Smirnov test of Normality:
ks.test(new_obs_model$residuals,"pnorm",0,sd(new_obs_model$residuals))
```

```
##
##  One-sample Kolmogorov-Smirnov test
##
## data:  new_obs_model$residuals
## D = 0.0672809, p-value = 0.99778
## alternative hypothesis: two-sided
```

```
### Anderson-Darling test of Normality:
require(nortest)
ad.test(new_obs_model$residuals)
```

```
##
##  Anderson-Darling normality test
```

```
##
## data:  new_obs_model$residuals
## A = 0.202652, p-value = 0.86542
### Jarque-Bera test of Normality:

require(tseries)
jarque.bera.test(new_obs_model$residuals)

##
##  Jarque Bera Test
##
## data:  new_obs_model$residuals
## X-squared = 0.653021, df = 2, p-value = 0.72144
```

The extremely high p-values for all of the above tests cause us to fail to reject the null hypothesis, which is that the residuals are Normal.

Constant variability :

```
require(olsrr)
ols_test_breusch_pagan(m1)

##
##  Breusch Pagan Test for Heteroskedasticity
##  -----
##  Ho: the variance is constant
##  Ha: the variance is not constant
##
##               Data
##  -----
##  Response : runs
##  Variables: fitted values of runs
##
##          Test Summary
##  -----
##  DF          =      1
##  Chi2         =    0.014290331
##  Prob > Chi2  =    0.90484583
```

The high p-value provides evidence in favor of the null hypothesis, i.e., the variance is constant.